What Offline Fitting Buys Under Benchmark Shift: Predicting LLM Item Outcomes from a Few Labels

Arshia Teymourtash*
KTH Royal Institute of Technology
Stockholm, Sweden
arturt@kth.se

Ankit Grover*
KTH Royal Institute of Technology
Stockholm, Sweden
agrover@kth.se

REVIEW COPY, not a paper version. Light red = existing text I would delete or rewrite, or (marked [S1 was] to [S20 was]) the old wording of an edit already made. **Light green** = text or an analysis I would add, or (marked [S1 now] to [S20 now]) the new wording of an edit already made. Numbers in brackets refer to the list *What I would change and why*; [S1] to [S20] to the list of edits already made. Both lists are at the end.

STILL TO IMPROVE [A-title]: Title: still undecided. Option on the table: a title that states the finding, e.g. "Offline Fitting Transfers a Prior, Not Item Estimates: Predicting LLM Item Outcomes on Unseen Benchmarks".

Abstract

STILL TO IMPROVE [A-abstract]: Abstract: still long and number-heavy. Tighten it to one claim per sentence, check every number against the tables, and settle the headline (the average over splits) before finalising.

Predicting how an LLM answers each item of a new benchmark from a few labelled answers is useful only if what is learned from past benchmarks carries over. We test what carries over with one controlled design: the same models, tasks and label order, changing only where the prior comes from (fixed in advance, tuned on validation, or fitted on training benchmarks), in an online simulation of the competition protocol with three unseen test benchmarks. On benchmarks seen in training, keeping offline estimates of item difficulty gains 0.085 Brier ALC over a Jeffreys Beta-Binomial baseline, nine times the gain of any prior alone. On unseen benchmarks that gain disappears, and only a shrinkage prior transfers, with a modest and uneven gain: 0.005 averaged over all 17 benchmarks, and 0.004 averaged over all 680 choices of three held-out benchmarks, where it hurts in 17.5% of them. Our test split is a favourable one: fitted on training, the prior gains 0.014 on test but loses 0.008 on validation. Across the 17 benchmarks, the gain tracks how far a benchmark's accuracy lies from the prior mean (Spearman 0.96). Gains concentrate at zero to three labels. The two neural probes we trained on item embeddings recover a small rank signal for item difficulty on the six unseen benchmarks (Spearman 0.09 and 0.10 against 0.03 for a random baseline), but using it as the prior mean of item difficulty does not improve Brier ALC. A well-fitted Beta-Binomial beats the Rasch models on test but not on validation; whether sharing labels across subjects helps is untested, since each draw has one subject per benchmark. Our final Codabench submission (1005155) is a Beta-Binomial predictor with a fixed prior and subject-level shrinkage. We release the simulation, the frozen draws and every saved score.

*Work done equally.

40th Conference on Neural Information Processing Systems (NeurIPS 2026). Track on Evaluations and Datasets.

1 Introduction

Evaluating language models on every item of every benchmark is expensive, so a predictor that infers a model's per-item results from a handful of labelled answers is attractive. The natural recipe is to learn from past benchmarks how items behave and carry that to a new benchmark. Whether it carries is an empirical question, and the answer matters because new benchmarks and new models keep arriving. Learned predictors are known to fail on unfamiliar models while estimators anchored to the observed labels survive [12], and amortized item-response models learn item difficulty from item text [11]. We ask what, among the things one can fit offline, actually transfers to an unseen benchmark.

We use the Predictive AI Evaluation Competition's setting: a *task* is one subject (a model under given evaluation settings) answering one benchmark, labels are bought one at a time, and predictions are scored by Brier ALC at 0, 1, 3, 7, 15 and 31 labels. Our contributions:

- **A controlled comparison.** Beta-Binomial, Rasch and kernel Rasch predictors share tasks, labels and order; only the *prior source* changes (Section 4).
- **A seen-benchmark check.** Offline fitting must win on benchmarks it has seen; it does, by a wide margin, and the margin vanishes on unseen benchmarks (Section 5).
- **When a transferred prior helps.** A prior fitted on training benchmarks helps on benchmarks whose accuracy is near the prior mean and hurts far from it, shown on 17 benchmarks, the 11 training ones by leave-one-benchmark-out.
- **Item features.** Neural probes on item embeddings carry a small rank signal for difficulty on unseen benchmarks that does not improve prediction as a prior mean (Section 5.4).
- **NEW [5]:** *A bullet for the negative results the reviewer asked about: adaptive priors (mixtures, strength averaging, a pilot switch rule) did not remove the harm without losing the gain, and the first labels warn only weakly (pooled AUC 0.73, within-benchmark 0.64).*
- **Open materials.** The frozen draws, saved scores and code to rerun every table and figure.

2 Related work

Empirical Bayes estimates a prior from related observations; classical shrinkage and small-area estimation show how pooling can stabilize noisy group estimates [1–3]. Our Beta-Binomial baseline is its simplest instance: a Beta prior on a task's accuracy, fixed in advance (Jeffreys [9]), tuned on validation, or fitted across training tasks. Prediction-powered inference also combines predictions with observed labels, but targets uncertainty about population quantities rather than sequential item-level prediction [4].

Item response theory explains a response by the subject's ability and the item's difficulty; the Rasch model is its basic form [10]. Extensions model item–respondent interactions in a latent space or separate language effects from item content [5, 6], and amortized IRT predicts item parameters from text [11]. Work on benchmark prediction from fewer data reports that learned predictors fail on unfamiliar models while estimators anchored to the observed labels survive [12]. We ask the analogous question for unfamiliar benchmarks: whether item-level signals and transferred priors remain useful on a new benchmark, where only a few labels are available.

3 Data and evaluation protocol

Data. The data are a binary response matrix $Y \in \{0,1\}^{N \times M}$ with N subjects (rows, index i) and M items (columns, index j): $Y_{ij} = 1$ if subject i answers item j correctly. Most entries are missing, since a subject answers only the items of the benchmarks it was evaluated on. Each item j belongs to the dataset $d(j)$ of one of $D = 17$ benchmarks, so a benchmark is a block of columns, and a *task* is one row restricted to one block: subject i on benchmark d.

Each item j has an embedding $x_j \in \mathbb{R}^{1024}$ computed once from its text by the frozen BGE-M3 encoder [8] with a context of 8,192 tokens, and normalised to unit length. We use BGE-M3 because

the corpus is multilingual (it includes Chinese, Russian and several African languages) and some items are long, which English-only and 512-token encoders handle poorly.

Online simulation and splits. The competition protocol is *online*: for each task, labels are revealed one at a time and the predictor updates from them. Whatever the predictor knows before the first label is *offline*, that is, fitted on the training benchmarks beforehand. Our *online simulation* runs this protocol on recorded responses, with disjoint column blocks for training, validation, and test (11, 3, and 3 benchmarks): priors are fitted offline on training, and validation selects settings. The validation and test benchmarks are *out of distribution* (OOD), since nothing from them is seen during fitting.

A *competition-shaped draw* samples one subject per benchmark, so a predictor cannot borrow labels from other subjects on the same benchmark; this simulates the competition protocol on public data. As a separate *in-distribution* (ID) evaluation, we hold out rows (subjects) within the training blocks, so its benchmarks and items have been seen during fitting, and retain the questions covered by the fit subjects. The offline comparator reuses its fitted benchmark rates and item difficulties where available, and its online update changes only the subject ability.

Budgets and scoring. For each subject–benchmark pair we freeze a split of distinct items into acquisition and evaluation pools, assigning the extra item to acquisition when the count is odd. A frozen random acquisition order supplies nested prefixes of 0, 1, 3, 7, 15, and 31 labels. At each budget the predictor conditions on the allowed prefix and predicts the same evaluation items.

Evaluation outcomes are used only to score those predictions. We average Brier scores over query items, then average repeated appearances within each task and give tasks equal weight. Brier ALC combines the six budgets with weights $(0.1, 0.2, 0.2, 0.2, 0.2, 0.1)$. Paired intervals resample whole tasks with their matched predictions.

4 Methodology

4.1 Models

Beta-Binomial. Let p be a task's accuracy, with a Beta prior $p \sim \mathrm{Beta}(\alpha, \beta)$. After n revealed labels, k of them correct, the posterior is $\mathrm{Beta}(\alpha + k, \beta + n - k)$ and the prediction, the same for every item of the task, is its mean

$$\hat{p} = \frac{\alpha + k}{\alpha + \beta + n}. \tag{1}$$

The prior acts like $\alpha + \beta$ extra labels, α of them correct. The Jeffreys prior is $\alpha = \beta = \frac{1}{2}$ [9]. This is the exact Bayesian sequential update; nothing is learned.

Rasch. The probability that subject i answers item j correctly is

$$p_{ij} = \sigma(\theta_i - b_j), \qquad b_j = \delta_j - \mu_{d(j)}, \tag{2}$$

where σ is the logistic function, θ_i is the subject's ability, and b_j is the item's difficulty. This is the Rasch model [10] with the item difficulty split into a benchmark part and an item part: δ_j is the item's deviation from the average difficulty $-\mu_d$ of its benchmark $d = d(j)$, so a larger benchmark rate μ_d means an easier benchmark. The priors are $\mu_d \sim \mathcal{N}(m_0, \sigma_\mu^2)$, $\theta_i \sim \mathcal{N}(0, \sigma_\theta^2)$ and $\delta_j \sim \mathcal{N}(0, \sigma_\delta^2)$. At every budget we refit the maximum a posteriori (MAP) estimate from scratch on all labels revealed so far in the draw; an unlabelled item gets $\delta_j = 0$.

Kernel Rasch. Kernel Rasch keeps Equation (2) and changes only the prior on the item deviations $\boldsymbol{\delta} = (\delta_1, \ldots, \delta_M)^\top$. Rasch treats them as independent Gaussians, $\delta_j \sim \mathcal{N}(0, \sigma_\delta^2)$; kernel Rasch makes them jointly Gaussian, with covariance given by the similarity of the item embeddings:

$$\boldsymbol{\delta} \sim \mathcal{N}\big(\mathbf{0},\, K\big), \qquad K = \sigma_\delta^2\big[(1-\rho)\, I + \rho\, S\big], \quad S_{jj'} = \exp\Big(\frac{\cos(x_j, x_{j'}) - 1}{T}\Big), \tag{3}$$

where x_j are the item embeddings of Section 3, $\rho \in [0, 1]$ weights the similarity part S and T is its temperature. At $\rho = 0$ the two priors coincide; for $\rho > 0$, similar items share difficulty, while items

of different benchmarks stay independent. After the refit, a labelled item keeps its fitted deviation, and an unlabelled item j receives the posterior mean given the labelled items L of its benchmark, $\hat{\delta}_j = \rho\,\sigma_\delta^2\, S_{jL} K_{LL}^{-1} \hat{\boldsymbol{\delta}}_L$, which is 0 at $\rho = 0$ as in Rasch. With one subject per benchmark, as in the competition-shaped draws, no labels are shared across subjects: the joint fit equals independent per-task fits (maximum difference 10^{-14} in our check), so Rasch there differs from the Beta-Binomial only in its prior.

NEW [5]: *A short subsection on adaptive priors, all closed form: Beta mixtures fitted by EM whose weights update by marginal likelihood on the revealed labels; averaging over the prior parameters (α, β); and a switch to Jeffreys when the first k labels sit far from the prior mean, with the threshold chosen on validation only.*

NEW [7]: *Convergence paragraph: the online Newton solver converged in all 6,000 audited fits (at most 7 steps); the offline fit first stopped at an iteration limit, and a tight refit moved the hyperparameters by at most 1.2% and no result by more than 0.0001.*

Offline model that keeps its fitted values. For the ID evaluation, the offline Rasch keeps its fitted μ_d and δ_j for a benchmark and items it has seen and updates only the subject ability θ_i online (a one-dimensional MAP on the task's own labels). For an unseen benchmark or item it falls back to m_0 and 0, i.e. to the fitted prior alone.

4.2 Probes

A probe predicts the difficulty of an item from its features alone, before any label of its benchmark is seen. The target is the mean binary correctness per item, centered within benchmark for fitting. We compare a ridge probe of item embeddings, a probe using log character count, and a sampled random score. We also train two small neural probes (MLPs) on the item embeddings, one on item pass rate and one on the fitted Rasch item difficulty, using 8,462 items from the 11 training benchmarks only.

We score a probe on each benchmark d by two numbers, and average each over benchmarks. Let y_j be the observed difficulty of item j and $\hat{y}_j$ the probe's prediction. The *Spearman correlation* $\rho_d = \mathrm{corr}\left(\mathrm{rank}(\hat{y}_j), \mathrm{rank}(y_j)\right)$ compares the two orderings of the items: it is 1 if the probe orders them perfectly, 0 if its order is unrelated to the truth, and -1 if it is reversed. The *pairwise accuracy* $A_d = \frac{1}{|P_d|} \sum_{(j,k)\in P_d} s_{jk}$ is the share of item pairs the probe orders correctly, where P_d is the set of pairs with $y_j \neq y_k$ (observed ties are excluded), $s_{jk} = 1$ if $\hat{y}_j - \hat{y}_k$ has the same sign as $y_j - y_k$, $s_{jk} = \frac{1}{2}$ if the probe predicts a tie, and $s_{jk} = 0$ otherwise. A random score has $A_d = 0.5$ on average. The ID diagnostic holds out items within the 11 training benchmarks; the OOD columns evaluate entire unseen benchmarks. All methods use matched items and sampled pairs.

To test whether a probe helps prediction, we give the item deviations of Rasch a prior mean taken from it,

$$\delta_j \sim \mathcal{N}\left(\gamma\, z_j,\ \sigma_\delta^2\right), \tag{4}$$

where z_j is the probe's predicted difficulty of item j, centred within its benchmark (the benchmark rate μ_d absorbs the average), and γ is a scale; $\gamma = 0$ is the fitted Rasch. The probes use item features only, never outcomes.

STILL TO IMPROVE [A-adaptive]: Methodology: adaptive priors are missing. Add the mixture prior, strength averaging and the pilot switch rule as models, with their equations (see NEW [5] above), before the results on them can be reported.

4.3 Priors

STILL TO IMPROVE [A-priors]: Priors (Section 4.3): the bullets are dense. Separate prior parameters from prior sources more clearly, consider a small table of the settings per model and source, and state plainly that the fitted kernel Rasch keeps ρ and T from validation.

Each model has prior parameters: (α, β) for the Beta-Binomial; m_0 and the widths $\sigma_\mu, \sigma_\theta, \sigma_\delta$ for Rasch; and in addition ρ and T for kernel Rasch. Each is set from one of three sources.

- **Fixed in advance:** nothing learned. Beta-Binomial Jeffreys; Rasch $m_0 = 0$ and all widths 1; kernel Rasch additionally $\rho = 0.5$, $T = 0.1$.

- **Tuned on validation:** the best setting on the validation benchmarks from a grid (Beta-Binomial: $\alpha = 0.4$, $\beta = 0.6$; Rasch: $m_0 = 0.048$, $\sigma_\mu = 2$, $\sigma_\theta = 0.5$, $\sigma_\delta = 1$; kernel Rasch: $m_0 = 0.048$, $\sigma_\mu = 2$, $\sigma_\theta = 1$, $\sigma_\delta = 1.4$, $\rho = 0.8$, $T = 0.2$; probe prior mean, Section 4.2: the scale γ from a grid).
- **Fitted on training:** the 11 training benchmarks only. The Beta-Binomial prior maximises the Beta-Binomial marginal likelihood of the training tasks, weighting benchmarks equally ($\alpha = 1.92$, $\beta = 1.76$: a prior mean of 0.522 worth $\alpha + \beta = 3.68$ labels). The Rasch priors come from a joint MAP fit to 4.34 million training responses (4,809 subjects, 109,987 items) with empirical-Bayes hyperparameters averaged per benchmark so that the benchmark with most subjects does not dominate: $m_0 = 0.175$, $\sigma_\mu = 0.75$, $\sigma_\theta = 1.38$, $\sigma_\delta = 1.49$. Kernel Rasch uses the same values and keeps $\rho = 0.8$ and $T = 0.2$ from the validation tuning, so its kernel is not fitted on training. The probes of Section 4.2 are trained on the training benchmarks' items only. Only these numbers transfer; the fitted μ_d, θ_i and δ_j belong to training benchmarks, subjects and items.

5 Results

Every difference is written Δ_{a-b}, the Brier ALC of model a minus that of model b, so negative favours a; paired 95% bootstrap intervals are over tasks. J is the Jeffreys Beta-Binomial; BB, R and KR are the Beta-Binomial, Rasch and kernel Rasch models, subscripted fix, tune or fit for the source of their prior, and R off is the offline Rasch that keeps its fitted values.

5.1 Prior source matters, and no source wins on every split

Table 1 compares the nine combinations.

[8] On test, a Beta-Binomial with a prior fitted on training is the best predictor ($\Delta_{\mathrm{BB\,fit}-\mathrm{J}} = -0.0138$ [-0.0158, -0.0116]), four times the gain of the validation-tuned Rasch, and better than any Rasch model ($\Delta_{\mathrm{R\,fit}-\mathrm{BB\,fit}} = +0.0019$ [$+0.0007$, $+0.0031$]).

The validation-tuned Beta-Binomial prior is *worse* than Jeffreys on test ($\Delta_{\mathrm{BB\,tune}-\mathrm{J}} = +0.0055$) though it is the best Beta-Binomial on validation. Every comparison between prior sources changes sign between test and validation (Appendix B); fresh draws change no result by more than 0.0007 (Appendix C).

Table 1: **A prior fitted on training helps on the test benchmarks and hurts on the validation benchmarks; a fixed prior with strong shrinkage does nearly as well on test.** Each cell is $\Delta_{m-\mathrm{J}}$ for the model m in the row, competition-shaped draws. ID: held-out subjects on the 11 training benchmarks, 114 tasks, where the training-fitted Beta-Binomial prior is refit without those subjects. OOD validation: 3 unseen benchmarks, 80 tasks. OOD test: 3 unseen benchmarks, 45 tasks. Bold: best of the nine models in each column; underlined: second best.

Model m	ID	OOD validation	OOD test
Fixed in advance			
Beta-Binomial (Jeffreys)	0 (reference)	0 (reference)	0 (reference)
Rasch	-0.0073 [-0.0091, -0.0055]	$+0.0031$ [-0.0005, $+0.0066$]	-0.0102 [-0.0118, -0.0085]
Kernel Rasch	-0.0072 [-0.0087, -0.0058]	$+0.0029$ [-0.0008, $+0.0065$]	-0.0105 [-0.0121, -0.0088]
Tuned on validation			
Beta-Binomial	$+0.0041$ [$+0.0027$, $+0.0055$]	**-0.0015** [-0.0040, $+0.0010$]	$+0.0055$ [$+0.0038$, $+0.0072$]
Rasch	-0.0024 [-0.0028, -0.0019]	$+0.0004$ [-0.0006, $+0.0014$]	-0.0032 [-0.0035, -0.0029]
Kernel Rasch	$+0.0001$ [-0.0004, $+0.0005$]	$\underline{-0.0003}$ [-0.0010, $+0.0004$]	-0.0024 [-0.0027, -0.0021]
Fitted on training			
Beta-Binomial	**-0.0090** [-0.0115, -0.0063]	$+0.0079$ [$+0.0024$, $+0.0136$]	**-0.0138** [-0.0158, -0.0116]
Rasch	$\underline{-0.0081}$ [-0.0101, -0.0060]	$+0.0052$ [$+0.0008$, $+0.0097$]	$\underline{-0.0119}$ [-0.0133, -0.0106]
Kernel Rasch	-0.0063 [-0.0076, -0.0049]	$+0.0036$ [-0.0007, $+0.0076$]	-0.0118 [-0.0131, -0.0105]

NEW [9]: *One sentence and a pointer: log loss and calibration error give the same ordering on test and the same reversal on validation (log loss ALC* 0.6656 *against* 0.7045 *for Jeffreys on test;* 0.5404 *against* 0.5159 *on validation), and rank metrics are uninformative by construction for a constant prediction.*

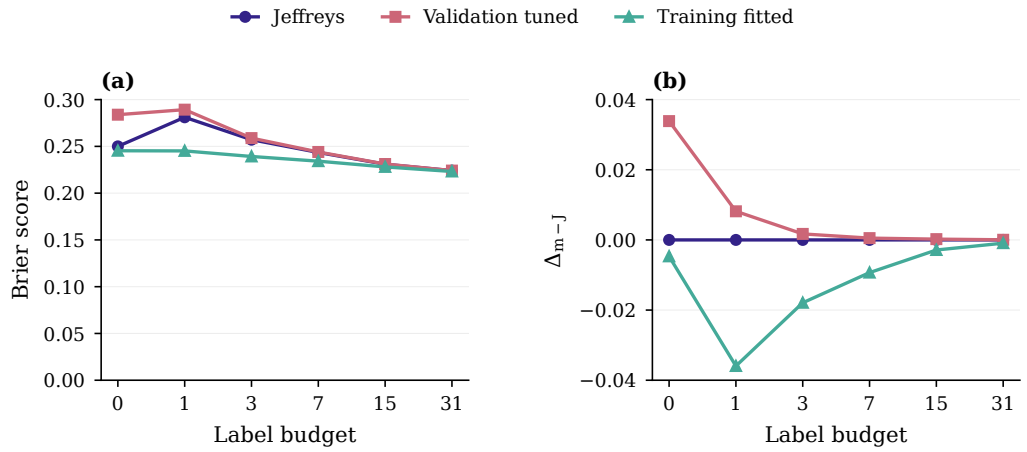


Figure 1: **The prior's source matters at zero to three labels and has faded by 31.** Brier score (a) and $\Delta_{m-\mathrm{J}}$ (b) per label budget on the three unseen test benchmarks, competition-shaped draws, for the Beta-Binomial with a fixed (Jeffreys, solid), validation-tuned (dashed) and training-fitted (dotted) prior. Jeffreys jumps to 0.75 or 0.25 after one answer, so the gap is largest at one label; at 31 labels the three differ by less than 0.001.

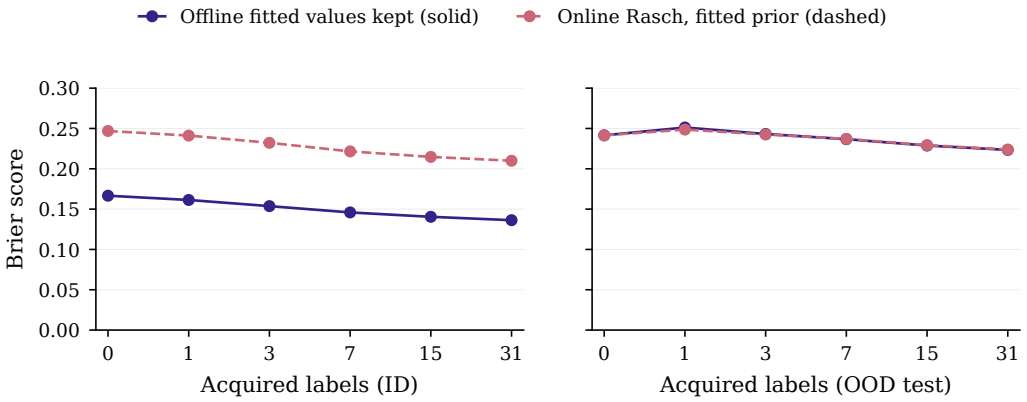


Figure 2: **Offline estimates win on familiar benchmarks and have no edge on unfamiliar ones.** Brier score per label budget for the offline Rasch keeping its fitted benchmark rates and item difficulties and updating only ability (solid), against the online Rasch with the same fitted prior (dashed). Left: ID, held-out subjects on the 11 training benchmarks. Right: OOD test benchmarks, where the offline model has no fitted values and falls back to the prior.

5.2 Learned priors transfer to unseen benchmarks; item-level estimates do not

On the ID evaluation set the offline model that keeps its fitted values scores $\Delta_{\mathrm{R\,off-J}} = -0.0852$ $[-0.0955, -0.0747]$ against $\Delta_{\mathrm{BB\,fit-J}} = -0.0090$ for the fitted Beta-Binomial prior and $\Delta_{\mathrm{R\,fit-J}} = -0.0081$ for the fitted Rasch prior (Figure 2); its item difficulties alone, at budget zero, reach a within-task AUC of 0.836. On the unseen test benchmarks it equals the fitted prior ($\Delta_{\mathrm{R\,off-J}} = -0.0115$ against $\Delta_{\mathrm{R\,fit-J}} = -0.0119$): nothing it knows about particular items or benchmark rates is available. This is the seen-benchmark check the design requires, and it separates two things that are easy to conflate: item-level estimates transfer only to the items they were fitted on, while the prior's spreads carry over to new benchmarks.

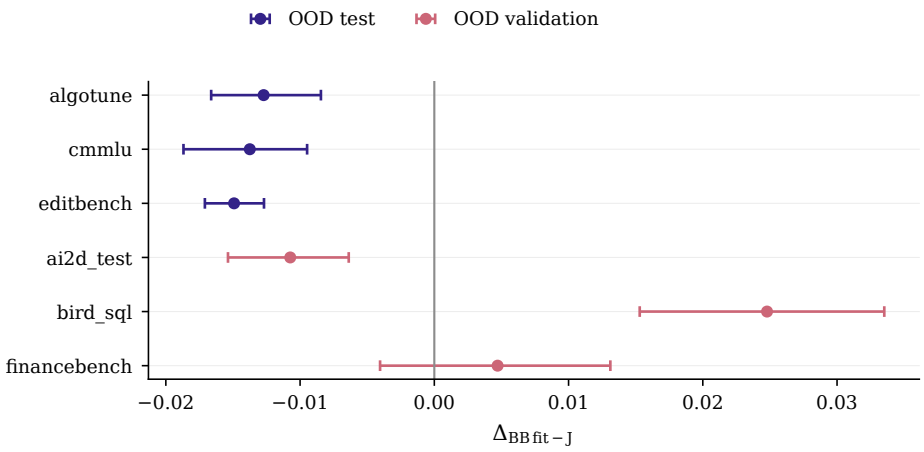


Figure 3: **The training-fitted prior helps on all three test benchmarks and on one validation benchmark, and hurts on the hardest.** $\Delta_{\mathrm{BB\,fit-J}}$ for each unseen benchmark, with 95% intervals over tasks. The validation benchmark `bird_sql` has task accuracy 0.12, far below every training benchmark.

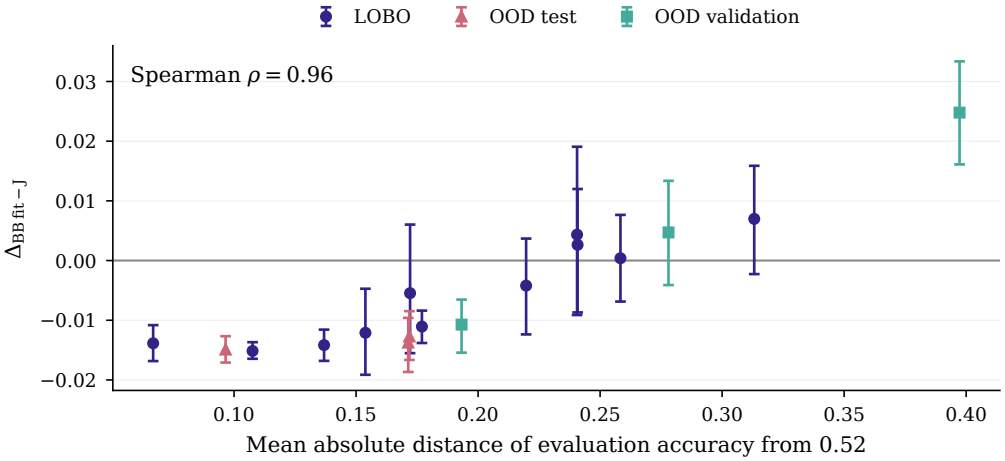


Figure 4: **A fitted prior helps when a benchmark's accuracy is near the prior mean.** Each point is one of 17 benchmarks. For each of the 11 training benchmarks, the prior was refit without any of that benchmark's rows and scored on held-out subjects. The six validation and test benchmarks use the original training-only prior. The horizontal axis uses task accuracy on the evaluation items; the vertical axis shows $\Delta_{\mathrm{BB\,fit-J}}$, with 95% task-bootstrap intervals. This is an explanatory analysis using observed outcomes.

5.3 Learned priors benefit benchmarks near the prior mean and hurt those far from it

QUESTION [Q1]: Section 5.3 and Figure 4: the text reported the 11-benchmark view (Spearman 0.93, helps on 5 of 11) while the Figure 4 caption already showed 17 benchmarks, and the abstract's Spearman matched the 11-benchmark text. Is the paper now consistent on 17 benchmarks everywhere (abstract, contributions, Section 5.3, Figure 4)?

Figure 3 shows the test and validation benchmarks separately: all three test benchmarks agree ($\Delta_{\mathrm{BB\,fit-J}}$ from -0.013 to -0.015), and on validation `ai2d_test` agrees while `bird_sql` goes

the other way ($+0.025$). Over the 17 benchmarks (Figure 4; each training benchmark is held out in turn, Appendix D) the fitted prior significantly helps on 9, significantly hurts on 1 (`bird_sql`), and the macro mean $\Delta_{\mathrm{BB\,fit-J}}$ is -0.0050. What separates the benchmarks is not whether they lie inside the range of training benchmark rates (`alignment_faking` and `harmbench` lose while inside it) but how far their task accuracies sit from the prior mean: the Spearman correlation between the gain and that distance is 0.96 over the 17 benchmarks (0.93 over the 11 training benchmarks), against 0.01 over the 11 for the benchmark's mean accuracy. Shrinking toward 0.52 helps tasks near 0.52 and hurts tasks far from it. The three test benchmarks (mean accuracy 0.54–0.68) sit near the prior mean, which is why test looks good. We found this pattern after seeing the results, so it is an explanation and not a predictor.

NEW [11]: *A short subsection,* How sensitive is the test gain to which benchmarks are held out?*, moved from the appendix: over 680 triples the mean gain is* -0.0037 *(prior refit on the other 14 benchmarks) or* -0.0051 *(training benchmarks only), the prior hurts in 17.5%–20% of splits, and our test triple gives* -0.0138 *when the prior is fitted on the 11 training benchmarks but* -0.0087 *when it is fitted on the other 14. State which fit set is used.*

NEW [12]: *The first-principles version of the explanation: for a task of true accuracy* p *the expected Brier excess of the posterior mean is* $\frac{(\alpha-p(\alpha+\beta))^2+n\,p(1-p)}{(\alpha+\beta+n)^2}$*, so the gain over Jeffreys is an exact function of* p *and of the budget; compare the predicted gain per benchmark with the observed one instead of citing a post hoc correlation.*

NEW [5]: *Results of the adaptive priors on the 17 benchmarks (worst harm, macro gain, test gain), and the pilot diagnostic: from 7 labels the distance of the pilot accuracy from the prior mean predicts harm with pooled AUC* 0.73 *(within a benchmark* 0.64*); the harm sits at 0–3 labels, before any pilot exists.*

Under the fitted priors, the simplest Rasch model is also the best on test and the worst on validation: with benchmark rate only $\Delta_{m-\mathrm{J}}$ is -0.0172 on test and $+0.0211$ on validation (Appendix A). The subject-ability term is what lets a benchmark far from the prior move away from it, at a cost in variance when the benchmark is near it. The kernel adds nothing on test.

5.4 Item content carries a small signal that does not improve the prior

Table 2: Item-difficulty transfer: within-benchmark Spearman correlation and pairwise accuracy between predicted and observed item difficulty, averaged across benchmarks. ID: items held out within the 11 training benchmarks; OOD validation and OOD test: three unseen benchmarks each. The MLP probes were not evaluated on ID (–).

	ID		OOD validation		OOD test	
Probe	Spearman	Pairwise	Spearman	Pairwise	Spearman	Pairwise
Random score	-0.003	0.499	-0.024	0.490	0.076	0.529
Log character count	0.107	0.539	0.058	0.525	0.061	0.521
Ridge on embeddings	0.400	0.649	0.196	0.570	-0.020	0.494
Pass-rate MLP	–	–	0.071	0.523	0.109	0.539
Rasch-target MLP	–	–	0.120	0.540	0.104	0.538

QUESTION [Q2]: Section 5.4: the probe table and the prior-mean table now sit in Results, with the probes defined in Section 4.2. Does the abstract say no more than these two tables show?

Embeddings show a clear signal on known benchmarks. On the six unseen benchmarks the two MLP probes pass the criterion we fixed in advance, with pooled Spearman 0.090 and 0.103 against 0.026 for a random baseline and a positive correlation on 5 of 6 benchmarks; the ridge probes do not (0.065 and 0.047). On the three test benchmarks alone the margin over the random baseline is small (0.109 and 0.104 against 0.076) and rests on three benchmarks.

This small signal does not help prediction. Table 3 shows that no scale γ of Equation (4) helps: every $\gamma \geq 0.25$ is worse than $\gamma = 0$ on both splits, and the best scale chosen on validation ($\gamma = 0.05$, Rasch-target MLP) is within 0.0004 of $\gamma = 0$ on both. This tests these two probes only; a stronger

difficulty predictor, for instance one trained through an item response likelihood on more items, could do better.

Table 3: Predicted item difficulty as the prior mean of the Rasch item deviation, Equation (4). Cells are $\Delta_{m-\mathrm{J}}$ of the fitted Rasch with scale γ, competition-shaped draws; $\gamma = 0$ is the fitted Rasch without a prior mean. Bold: lowest in each column.

	Pass-rate MLP		Rasch-target MLP	
Scale γ	OOD validation	OOD test	OOD validation	OOD test
0	$+\mathbf{0.0052}$	-0.0119	$+0.0052$	-0.0119
0.05	$+0.0062$	$\mathbf{-0.0119}$	$\mathbf{+0.0048}$	$\mathbf{-0.0123}$
0.1	$+0.0103$	-0.0082	$+0.0051$	-0.0120
0.25	$+0.0332$	$+0.0150$	$+0.0099$	-0.0079
0.5	$+0.0690$	$+0.0491$	$+0.0273$	$+0.0069$
1	$+0.1065$	$+0.0826$	$+0.0704$	$+0.0433$
1.5	$+0.1326$	$+0.1050$	$+0.1040$	$+0.0724$

6 Limitations

The filtered public data covers only three test benchmarks in the online simulation. Multiple draws and acquisition orders measure variation within this population; they do not supply independent tests of transfer to new benchmarks. Task bootstrap intervals likewise do not capture uncertainty from sampling new benchmarks. Results apply to the eligible response data and random label acquisition, and need not carry over to other benchmark mixes or selection policies. The 17-benchmark sensitivity analysis reuses held-out-subject tasks for the 11 original training benchmarks. The test triples overlap, so their distribution describes sensitivity within this dataset rather than independent estimates of transfer. The ID comparison has item response history that is unavailable under benchmark cold start. Observed item difficulty also depends on which subjects answered the item. The leave-one-benchmark-out association uses realized outcomes to explain when the fitted prior helps; it is not a way to select a prior before observing a new benchmark.

Further limits of our own. Each competition-shaped draw has one subject per benchmark, so sharing labels across subjects is not tested. The test gain depends on this split. The test benchmarks are close to the prior mean; over all 680 choices of three held-out benchmarks the fitted prior gains less ($\Delta_{\mathrm{BB\,fit}-\mathrm{J}} = -0.0037$ on average, refit on the other 14) and hurts in 17.5% of them (Appendix E), as it does on validation. Test scores were inspected repeatedly while building the models, so the test analyses are exploratory: all selection used validation, but test was not read once. The kernel settings were chosen on validation and not on the ID evaluation set. We did not rerun the subject-metadata prior in this design, and we dropped reasoning-trace features, which are not available at test time.

7 Conclusion

STILL TO IMPROVE [A-conclusion]: Conclusion: still to improve. The adaptive-priors sentence is unsupported until the body reports it; rewrite after the headline and the adaptive-priors result are settled, and keep it ranked and short.

What transfers from offline fitting to an unseen benchmark is a shrinkage prior, not item-level estimates. Where fitted values apply, offline fitting is worth a very large gain ($\Delta_{\mathrm{R\,off}-\mathrm{J}} = -0.0852$ on the ID evaluation set); where they do not, it is worth none. The prior's gain is modest and depends on the benchmark: -0.0050 on average over the 17 benchmarks and -0.0037 over the 680 choices of three held-out benchmarks, with a loss in 17.5% of them. The prior helps when a benchmark's accuracy lies near the prior mean and hurts when it lies far away, which is why the same fitted prior looked good on test and bad on validation. A well-fitted Beta-Binomial is the strongest estimator in our test, though not on validation; whether sharing labels across subjects helps remains untested, since each draw has one subject per benchmark. Two attempts to do better did not help: priors that adapt to the first labels did not remove the harm without losing the gain, and the neural probes we trained on item embeddings, which carry a small rank signal for difficulty, did not improve prediction as a prior mean.

Disclosure and rerun instructions

Training data. All fitted quantities (the offline Rasch fit, the Beta-Binomial empirical-Bayes prior, the item probes) use only the responses of the 11 training benchmarks in the competition's measurement database after the filtering described in Section 3; no other response data is used. **Pretrained models.** Item text is embedded with the frozen BGE-M3 model; nothing else is pretrained or fine-tuned. **LLM coding assistants.** Claude Code (Anthropic) assisted with writing code, analysis scripts and drafting text; every reported number was produced by running the code in the repository, and the authors checked the numbers. OpenAI Codex assisted with the reviewer analyses, figures, and text revisions.

The shared handoff contains frozen draws, scores, and competition-shaped test predictions. The GitHub analysis branch contains online-simulation and analysis code, aggregate results, and figure PDFs. Keep private response and prediction files in the authorized shared storage. The command `python organized/testing/eandd/wp9.py` with `-handoff` pointing to the local handoff and `-output` pointing to a fresh directory rebuilds the figures and checks the saved Brier scores against the saved predictions. The analysis README identifies the handoff location and documents how to render figures from aggregate CSVs alone. The WP7 item results are in `paper/eandd/data/wp7-summary.csv`. Record the analysis commit and retain the frozen draws for any rerun.

References

[1] H. Robbins. An empirical Bayes approach to statistics. In *Proceedings of the Third Berkeley Symposium on Mathematical Statistics and Probability*, 1:157–163, 1956.

[2] B. Efron and C. Morris. Stein's estimation rule and its competitors—an empirical Bayes approach. *Journal of the American Statistical Association*, 68(341):117–130, 1973.

[3] J. N. K. Rao and I. Molina. *Small Area Estimation*. Wiley, 2nd edition, 2015.

[4] A. N. Angelopoulos, S. Bates, C. Fannjiang, M. I. Jordan, and T. Zrnic. Prediction-powered inference. *Science*, 382(6671):669–674, 2023.

[5] M. Jeon, I. H. Jin, M. Schweinberger, and S. Baugh. Mapping unobserved item–respondent interactions: A latent space item response model with interaction map. *Psychometrika*, 86(2):378–403, 2021.

[6] G. Lior, T. Frostig, G. Stanovsky, and M. Eyal. Extending item response theory for efficient and meaningful multilingual evaluation. arXiv:2606.15643, 2026.

[7] G. W. Brier. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1), 1950.

[8] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, Z. Liu. BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation. arXiv:2402.03216, 2024.

[9] H. Jeffreys. An invariant form for the prior probability in estimation problems. *Proc. Royal Society A*, 186, 1946.

[10] G. Rasch. *Probabilistic Models for Some Intelligence and Attainment Tests*. Danish Institute for Educational Research, 1960.

[11] S. T. Truong, Y. Tu, P. Liang, B. Li, S. Koyejo. Reliable and efficient amortized model-based evaluation. In *ICML*, PMLR 267, 2025.

[12] G. Zhang, F. E. Dorner, M. Hardt. How benchmark prediction from fewer data misses the mark. arXiv:2506.07673, 2025.

A Rasch terms under the fitted priors

Table 4: $\Delta_{m-\mathrm{J}}$ for Rasch with terms switched off (prior width σ set near zero), under the training-fitted priors. Competition-shaped draws on OOD validation and OOD test. Bold: lowest in each column.

Model m	OOD validation	OOD test
Benchmark rate μ_d only	$+0.0211$ $[+0.0114, +0.0311]$	-0.0172 $[-0.0202, -0.0142]$
+ subject ability θ_i	$\mathbf{+0.0029}$ $[-0.0003, +0.0060]$	-0.0089 $[-0.0100, -0.0080]$
+ item deviation δ_j (Rasch)	$+0.0052$ $[+0.0007, +0.0098]$	-0.0119 $[-0.0133, -0.0106]$
+ kernel	$+0.0036$ $[-0.0006, +0.0078]$	-0.0118 $[-0.0131, -0.0105]$
μ_d and δ_j, no θ_i	$+0.0282$ $[+0.0164, +0.0401]$	$\mathbf{-0.0174}$ $[-0.0208, -0.0139]$
θ_i and δ_j, no μ_d	$+0.0078$ $[+0.0022, +0.0135]$	-0.0140 $[-0.0157, -0.0122]$

B Paired comparisons between prior sources

Table 5: Paired Δ_{a-b} between prior sources, with 95% intervals over tasks. Each comparison changes sign between OOD test and OOD validation. A dash: not run on ID.

Comparison	ID	OOD validation	OOD test
$\Delta_{\mathrm{R\,fit}-\mathrm{R\,fix}}$	–	$+0.0021$ $[+0.0006, +0.0036]$	-0.0017 $[-0.0028, -0.0005]$
$\Delta_{\mathrm{KR\,fit}-\mathrm{R\,fit}}$	–	-0.0016 $[-0.0021, -0.0012]$	$+0.0001$ $[-0.0002, +0.0004]$
$\Delta_{\mathrm{R\,fit}-\mathrm{BB\,fit}}$	$+0.0009$ $[-0.0001, +0.0018]$	-0.0027 $[-0.0041, -0.0012]$	$+0.0019$ $[+0.0007, +0.0031]$

C Draw-seed robustness

Table 6: $\Delta_{m-\mathrm{J}}$ on three fresh sets of 150 competition-shaped draws per role (seeds 20261201–03); results on the frozen draws are in Table 1.

Model m	OOD validation	OOD test
BB fit	$+0.0078$ / $+0.0083$ / $+0.0089$	-0.0137 / -0.0144 / -0.0139
R fit	$+0.0051$ / $+0.0054$ / $+0.0060$	-0.0120 / -0.0124 / -0.0123
BB tune	-0.0015 / -0.0016 / -0.0018	$+0.0060$ / $+0.0059$ / $+0.0063$

D Leave-one-benchmark-out for the fitted prior

NEW [16]: *Extend this table to all 17 benchmarks (the figure already shows 17), with the Spearman correlation for 17 (0.96) beside the one for the 11 training benchmarks (0.93).*

Table 7: $\Delta_{\mathrm{BB\,fit-J}}$ with the prior refit without the held-out benchmark, on each training benchmark's held-out-subject tasks, with the mean distance of its task accuracies from the prior mean 0.52.

Held-out benchmark	Tasks	$\Delta_{\mathrm{BB\,fit-J}}$ [95% CI]	Mean distance from 0.52
`hallusionbench`	20	-0.0151 [-0.0165, -0.0136]	0.11
`bridging_gap`	20	-0.0142 [-0.0167, -0.0114]	0.14
`helm_afr`	3	-0.0139 [-0.0168, -0.0108]	0.07
`afrimedqa`	6	-0.0121 [-0.0189, -0.0046]	0.15
`afrieval`	20	-0.0111 [-0.0138, -0.0085]	0.18
`cruxeval`	5	-0.0055 [-0.0155, $+0.0056$]	0.17
`brace`	6	-0.0042 [-0.0121, $+0.0037$]	0.22
`agi_elo`	4	$+0.0004$ [-0.0069, $+0.0076$]	0.26
`appworld`	4	$+0.0026$ [-0.0087, $+0.0120$]	0.24
`alignment_faking`	20	$+0.0043$ [-0.0092, $+0.0188$]	0.24
`harmbench`	6	$+0.0070$ [-0.0023, $+0.0161$]	0.31
Mean over the 11 benchmarks	114	-0.0056	0.19

NEW [7]: *New appendix tables: the convergence audit; the adaptive-prior results per benchmark; the pilot-diagnostic AUC and the switch-rule sweep; and the per-split distribution for the Rasch prior (40 triples).*

E Sensitivity to the benchmark split

We held out each of the 680 triples, refitted the benchmark-weighted empirical Beta prior on the other 14 benchmarks, and rescored the frozen trajectories. The mean macro $\Delta_{\mathrm{BB\,fit-J}}$ was -0.0037, with 5th–95th percentiles $[-0.0095, +0.0099]$. The fitted prior hurt on 17.5% of splits. The original test benchmark names gave -0.0087 (percentile 13) under refitting. This refit includes the former validation benchmarks among the 14 eligible fit benchmarks for the original test triple, so its result differs from the original 11-benchmark fit. The 11 original training benchmarks are scored on held-out-subject tasks from the existing leave-one-out data. These overlapping triples describe split sensitivity within this dataset.

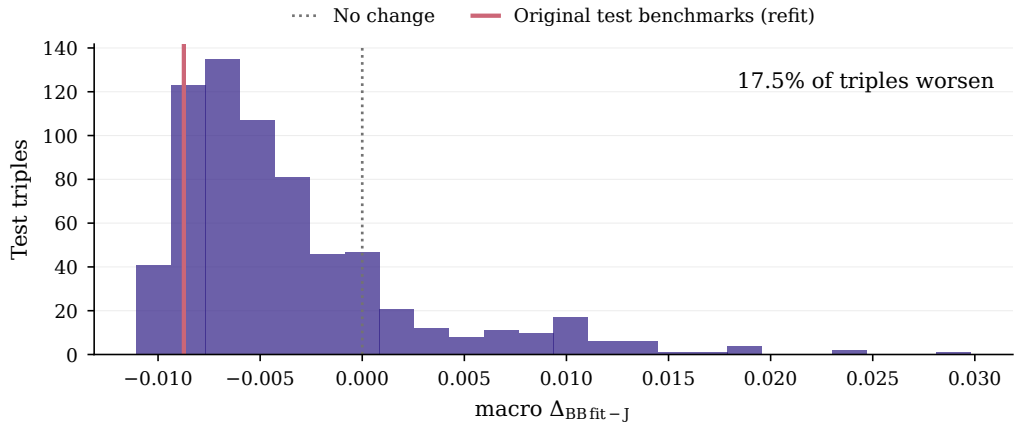


Figure 5: $\Delta_{\mathrm{BB\,fit-J}}$ over all 680 triples. Each empirical prior is fitted on the 14 benchmarks outside its triple. The vertical line marks the original test benchmark names under this refitting protocol.

F Alternative evaluation metrics

Table 8: Competition-shaped online simulation; models are labelled as in Section 5. Brier, log loss, and ten-bin task ECE are averaged using the official budget weights; lower is better. AUC uses within-task item rankings at 31 labels, where higher is better. Dashes indicate no informative ranking; a constant prediction ties all items. Bold: best in each split (lowest loss, highest AUC). Per-budget AUC and alternative weights are reported in the accompanying CSVs.

Split	Model m	Brier	Log loss	ECE	AUC at 31 labels
ID	J	0.2357	0.6681	0.1407	–
	BB fit	0.2268	0.6437	0.1235	–
	R fit	0.2276	0.6460	0.1245	–
	KR fit	0.2295	0.6506	0.1336	0.581
	R off	**0.1506**	**0.4634**	**0.1161**	**0.836**
OOD validation	J	**0.1738**	**0.5159**	**0.1405**	–
	BB fit	0.1817	0.5404	0.1677	–
	R fit	0.1790	0.5323	0.1580	–
	KR fit	0.1774	0.5279	0.1603	**0.618**
OOD test	J	0.2499	0.7045	0.1426	–
	BB fit	**0.2362**	**0.6656**	**0.1113**	–
	R fit	0.2380	0.6712	0.1164	–
	KR fit	0.2382	0.6717	0.1230	**0.571**

What I would change and why

1. **Headline.** 0.014 is the best case: our test split is one of the most favourable we could have drawn. The average over splits is about 0.004–0.005, and the prior hurts in roughly one split in five.
2. **17 benchmarks.** Figure 4 already shows 17 benchmarks (Spearman 0.96, helps on 9, hurts on 1); the text, the abstract and the appendix table still say 11 (Spearman 0.93, 5 of 11).
3. **Item features.** A pass-rate MLP and a Rasch-target MLP pass the rank criterion set in advance (pooled 0.090 and 0.112 against 0.026 for a random control); ridge probes do not. On the three test benchmarks alone the margin over the control is small (0.109 against 0.076). The claim of no signal must be softened; the Brier test of the prior-mean option decides whether it matters.
4. **Label sharing.** "Adds nothing to a good prior" holds on test and with one subject per benchmark. On validation the fitted Rasch beats the fitted Beta-Binomial, and multi-subject draws are barely tested.
5. **Negative results the reviewer asked for.** Adaptive priors and the pilot switch rule do not meet the criteria fixed in advance; the first labels warn only weakly.
6. **Specificity.** Say which widths and kernel settings the validation-tuned Rasch priors took.
7. **Convergence and extra appendix tables** (answers a highlighted reviewer point).
8. **Wording of the best result.** Add that it is the favourable split and that the validation split reverses.
9. **Other metrics.** Log loss and calibration agree with Brier (reviewer question 3).
10. **Correct the 11-benchmark numbers** in the running text to the 17-benchmark view.
11. **Sensitivity to the split in the main text,** with the fit set stated: -0.0138 (11 training benchmarks) against -0.0087 (the other 14).
12. **An exact risk formula** so the explanation is a derivation, not a correlation found after the fact.
13. **Limitations.** Quantify the favourable-split statement with the 680 splits.
14. **Conclusion.** Rewrite around the average, the spread across splits and the limits of what a label-based diagnostic can do.
15. **Terminology.** Done: the table caption now says "online simulation" (see S5).
16. **Appendix table** to 17 benchmarks.

Edits already made to the paper (old wording red, new wording green)

S1 Beta-Binomial notation. Conventional symbols (checked against standard references): success probability p, prior $\mathrm{Beta}(\alpha, \beta)$, n trials with k successes, posterior $\mathrm{Beta}(\alpha + k, \beta + n - k)$ and posterior mean $(\alpha + k)/(\alpha + \beta + n)$, with $\alpha + \beta$ read as the number of pseudo-labels the prior is worth. The old text used π, a mean m and a strength s (which clashed with the subject index s) and h for the correct labels. The prior settings are now stated as (α, β): Jeffreys $(\frac{1}{2}, \frac{1}{2})$, validation-tuned $(0.4, 0.6)$, training-fitted $(1.92, 1.76)$.

S2 Rasch equation in IRT convention. $p_{ij} = \sigma(\theta_i - b_j)$, replacing $\mathrm{logit}\, p = \mu_b + \theta_s - b_i$. The item difficulty is b_j (see S4 for its split); the IRT symbol β is left to the Beta prior.

S3 Response matrix. Section 2 now defines $Y \in \{0, 1\}^{N \times M}$ with rows i (subjects) and columns j (items), benchmarks as blocks of columns, a task as one row restricted to one block, and the splits as disjoint column blocks; the ID control holds out rows.

S4 Symbols. Item difficulty $b_j = \delta_j - \mu_{d(j)}$: δ_j is the deviation from the benchmark's average, μ_d the benchmark rate, $d(j)$ the dataset (benchmark) of item j, $D = 17$ benchmarks. Prior widths are standard deviations $\sigma_\mu, \sigma_\theta, \sigma_\delta$ (formerly τ); the kernel uses an indicator $\mathbf{1}_{j=j'}$.

S6 Explicit differences. Every difference is Δ_{a-b}, the Brier ALC of model a minus model b, with J (Jeffreys), BB, R and KR (Beta-Binomial, Rasch, kernel Rasch) and fix, tune, fit or off as the subscript, in the Results text, table headers, the paired and robustness tables and the captions. **Not yet done:** the axis labels inside Person-B's figure files still read "difference to Jeffreys" or "Fitted minus Jeffreys".

S7 Related work is a numbered section after the Introduction instead of a bold run-in paragraph.

S8 Title block. Authors, affiliation, address and emails added (were *Anonymous Author(s)* and the repository link). The page footer keeps the NeurIPS 2026 conference line.

S9 Methodology. Section 4, formerly *Models and prior sources*, is now *Methodology* with two subsections: *Models* (Beta-Binomial, Rasch, kernel Rasch as its own paragraph, the offline model) and *Priors* (a sentence listing each model's prior parameters, then the three sources: fixed, tuned, fitted).

S10 **Appendix.** Same symbols as the main text: rows of the Rasch-terms table use $\mu_d, \theta_i, \delta_j$ and the width σ; table rows and the paired, robustness and metrics tables use the model labels J, BB, R, KR (fix, tune, fit, off) with Δ_{a-b}; "panel" removed.

S11 **Section 3 structure.** The protocol text is split into bold run-in paragraphs: Data, Splits, Draws, Label acquisition, Scoring, In-distribution control. The wording is unchanged apart from S3.

S12 **Kernel Rasch.** The two item priors are shown side by side as equations: Rasch $\boldsymbol{\delta} \sim \mathcal{N}(\mathbf{0}, \sigma_\delta^2 I)$ and kernel Rasch $\boldsymbol{\delta} \sim \mathcal{N}(\mathbf{0}, K)$ with $K = \sigma_\delta^2[(1-\rho)I + \rho S]$, so the only difference is visible; they coincide at $\rho = 0$.

S13 **Abstract.** "language model" is now "LLM"; the rest of the abstract is unchanged, so its claims still await [1], [3] and [4].

S14 **Abstract, authors.** Added the 17-benchmark average and the 680-split average with the share of harmful splits (the favourable-split point of [1]); replaced the item-features sentence with the probe result and the R5 null; added the equal-contribution mark (asterisk on both names, footnote "Work done equally.").

S15 **Locked in.** Items [2], [3], [10], [13] and [14] are now in the paper and no longer shown as red or green: the 17-benchmark view (Spearman 0.96, helps on 9 of 17, hurts on 1, macro -0.0050) in the abstract, the contributions and Section 5.3; the two neural probes and the R5 null in Section 6; the 680-split sentence in the limitations; the rewritten conclusion.

Q1, Q2 **Open questions** on the 17-benchmark consistency and on Section 6 are shown as green question boxes in the text.

S16 **Embeddings and kernel settings.** A new *Item embeddings* paragraph (frozen BGE-M3, 1024 dimensions, 8,192-token context, unit length) is in Section 3 and Kernel Rasch refers to it; Section 3 is four bold paragraphs (Data, Item embeddings, Splits and draws, Label acquisition and scoring); the kernel paragraph states how unlabelled items are predicted; the settings moved to *Priors*: tuned values for Rasch and kernel Rasch, and that the fitted kernel Rasch keeps $\rho = 0.8$, $T = 0.2$ from validation.

S17 **Honest claims.** Abstract and conclusion lead with the average over 17 benchmarks and over the 680 splits and the share of harmful splits, then give the test and validation gains as the two examples. The label-sharing claim is replaced by a statement that it is untested (one subject per benchmark in every draw), also added to the limitations. The item-feature claim is limited to the two probes we trained (Section 6 notes a stronger predictor could do better).

S18 **Item-difficulty section dissolved.** The separate Section 6 is gone. *Probes* is a new subsection of Methodology (4.2: the probes, how they are scored, and Equation 5, the probe prior mean); the two tables and the discussion are Results 5.4, *Do item features help?*; the scale γ and the training-only probes are noted in *Priors*; a contribution bullet was added.

S19 **Research questions and headings.** The Results subsections are statements of what each shows (no separate list of research questions: the Introduction keeps its one governing question and the contribution bullets). Table headers read ID, OOD validation, OOD test (Table 1 reordered to match; Table 2 caption gives the benchmark counts).

F1, F3, F4, F5 **Figure labels: done by Person-B (commit 2bea768).** The axis labels and legends inside Figures 1, 3, 4 and 5 still use the old wording ("difference to Jeffreys", "Fitted minus Jeffreys", "Test", "Validation"). They are drawn by Person-B's plot scripts, so Person-B must regenerate them with Δ_{a-b} labels and OOD legends; shown as red boxes at each figure.

S20 **Tables.** All tables follow NeurIPS style (booktabs, captions above, no vertical rules). Table 1 and Tables 3, 4 and 8 bold the best value per column with a note in the caption; Table 1 shows the reference as "0 (reference)"; Table 2 has grouped columns (Spearman, Pairwise under ID, OOD validation, OOD test) with the random score as a baseline row; Table 4 separates the added terms from the ablations; Table 7 has a mean row; Table 8 has a rule between splits and the split label on the first row of each block.

S5 **Wording.** Section 2 is titled *Data and evaluation protocol* (was *Setting and online simulation*); "panel" became "ID control set" and "response data"; "replay" became "online simulation".